

Architectural Foundations for Embedded Systems

Ved Shanbhogue, Rivos Inc.

Use Cases



- Wide array of embedded usages - industrial, telecom, medical, instrumentation, test and measurement, aerospace, automotive, retail, etc.
- Varying power and performance requirements
- Enabling these usages need some/all of these Foundational technologies
 - Security, Reliability, Deterministic performance, Isolation, confidential computing, telemetry

Foundational Technologies

Confidential Computing

Virtualization

Quality Of Service

Reliability, Availability, Serviceability

Secure Software


Secure and Resilient Platforms

- Detect/Prevent/Contain exploitation of SW/FW bugs
- Remote attacks, Physical attacks, persistent threats, supply chain attacks
- Maximize uptime and ensure correct operation of the system
- Error containment, ease of system integration, TCO
- Determinism, real time deadlines, bounded worst case execution time

Secure Resilient Platforms

- **Strong roots of trust for - ROT/COT - for update (RTU), detection of corruption (RTD), and recovery (RTRec)**
- **Boot Security** - Measured Boot
 - Protect system compromise during startup
 - Measure each layer of firmware & SW in chain of trust
 - Reliably record measurements for verification
 - Secure debug authorization
 - Attestation: evidence of identity, state, configurations, FW/SW images, etc.
 - Sealing: Encrypted data can only be decrypted by same device in same state
 - Standards and references
 - TCG Device Identifier Composition Engine (DICE)
 - Open Titan - Secure Boot
 - TPM based
- **Firmware resilience** - of platform firmware and critical data such as configurations
 - Protect: f/w Integrity; update authentication, resilient storage, anti-rollback
 - Detect: Corruption of f/w and critical data
 - Recover: From corruptions
 - Guidelines
 - NIST SP 800-193 - Platform Firmware Resiliency Guidelines
 - DMTF Redfish

Foundations of a secure resilient platform



Secure Resilient Platforms

- **Platform Attestation - secure reporting of platform ID and state to remote verifiers**

- Provisioning, Job scheduling, Fleet Auditing
- Supply chain security
- Support many roots of trusts in the platform
- Identity of devices and measurement of device firmware and configurations
- Published attestation policies - Expected measurements, latest versions
- Standards and references
 - DMTF security protocol and data model (SPDM)
 - PCI-SIG Component Measurement and Authentication (CMA)
 - USB type-C authentication specification
 - IETF Remote Attestation Procedures (RATS)
 - IETF Concise Reference Integrity Manifest (CoRIM, CoMID, CoSWID)

- **Data protection**

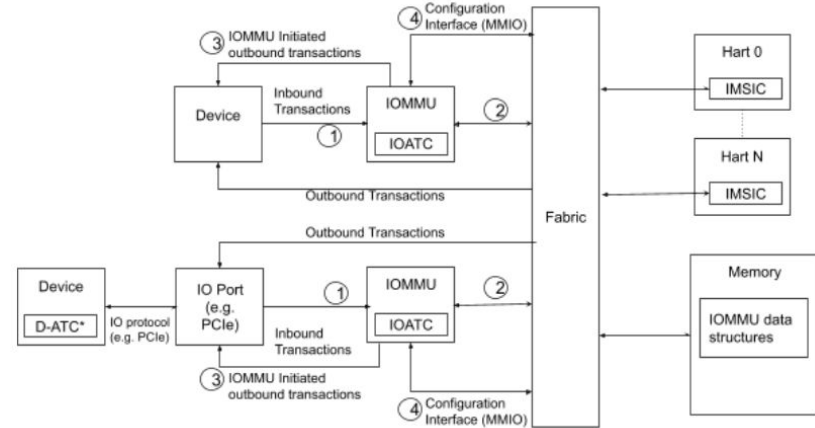
- Sealing: Encrypted data can only be decrypted by same device in same state
- Non-volatile storage encryption and integrity
- Memory encryption and integrity
- PCIe and CXL link Integrity and Data Encryption (IDE)

Secure Software

- Harden software against common and emerging attack techniques
- Available ISA capabilities
 - Privilege separation - M/S/U
 - Page Based memory protection
 - PMP/ePMP
 - Virtualization
 - Defend privilege escalations - executing from user-mode pages, accidentally consuming user-mode data
- Control-flow integrity (CFI) - **gap**
 - Commonly known as code-reuse attacks and one of most common exploit techniques
 - e.g., Return-oriented programming (ROP), Jump-oriented programming (JOP)
 - Divert control flow by overwriting control flow variables (return addresses, function pointers)
 - CFI-SIG formed to develop strategy for mitigation or prevention of these techniques
- Memory Safety Enforcement - **gap**
 - Memory safety bugs such as buffer overflows, heap-use-after-free, and stack-use-after-return.
 - Commonly exploited to cause remote code execution, data leakage, privilege escalation, etc.
- Transient execution
 - Transient instructions may perform unauthorized/unintended computations
 - Reveal unauthorized results - secrets - through side channels (Specter, Meltdown, etc.)
 - Mostly a function of secure hardware design
 - Need ISA constructs to control speculation (e.g. FENCE.T) - **gap**

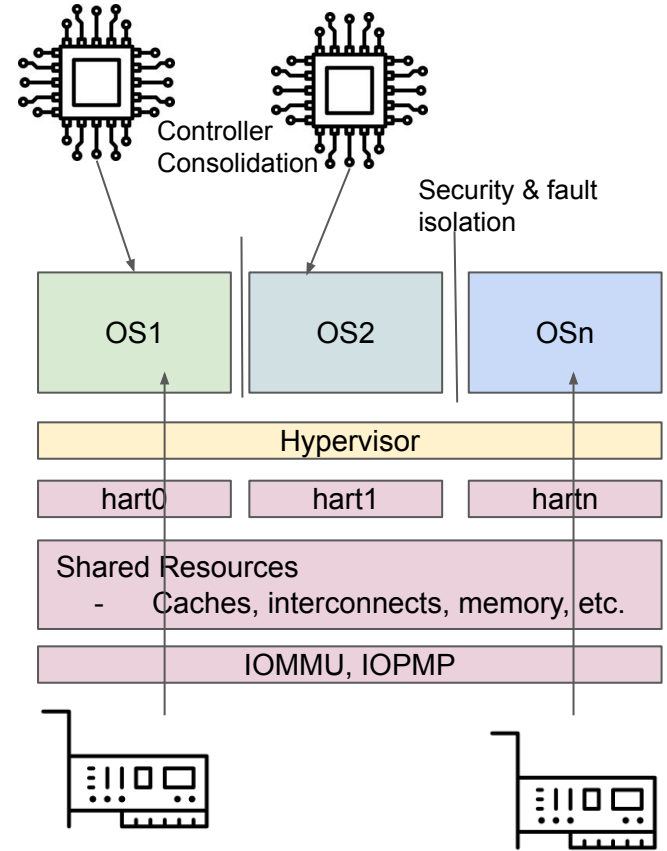
Secure Software

- IOMMU and IOPMP
 - Required even when virtualization not enabled
 - Defend against malicious DMA attacks
 - Defend against driver bugs
 - Protect cross VM attacks and attacks on hypervisor
 - Malicious/buggy VMs using direct assigned devices
 - Functions
 - IO virtual addresses => physical addresses
 - Page based protections to device DMA
 - Address Translation Services
 - Shared virtual memory with devices
 - Quality of Service Enforcement
 - Direct interrupt routing to guests



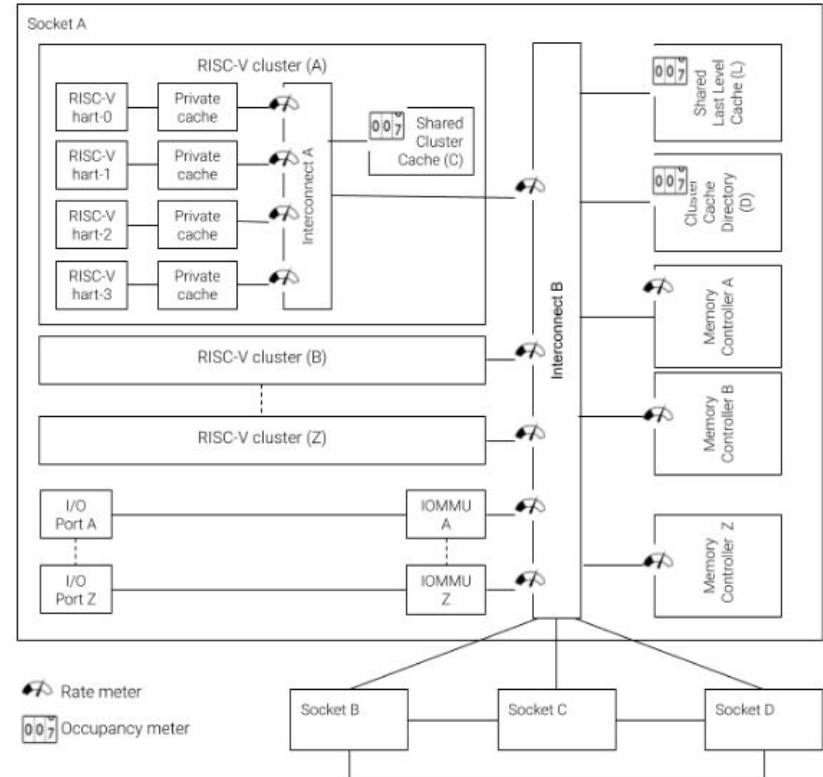
Virtualization

- Consolidate software on multiple embedded controllers to a single platform
- Fault isolation by executing critical applications in separate VMs
- Mixed criticality usages
 - Logically independent software (e.g. IVI, motor control) stacks in their own VMs
- Enhanced security through VM based isolation
- Ease of integration, updates, management of software stacks
- DMA protection and direct device assignment - needs IOMMU
- Challenges
 - Sharing of resources between multiple virtual machines
 - Caches, interconnects, memory controller, power, etc.
 - Lead to non-determinism
 - Need hardware mechanisms for controlling resource usage



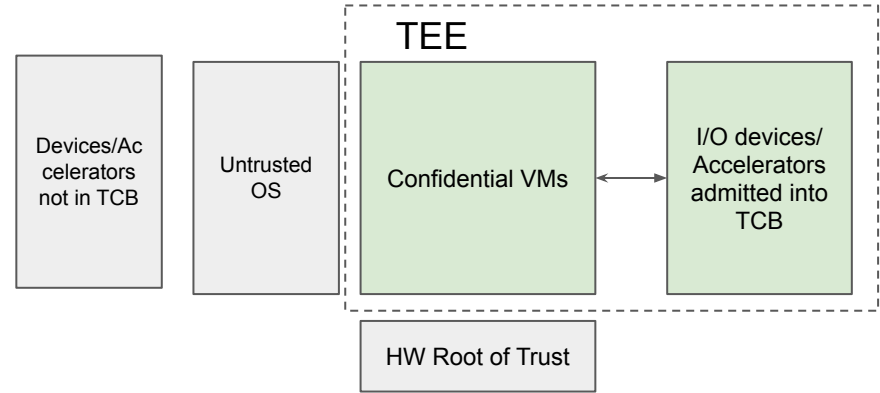
Quality Of Service

- Performance guaranteed in advance by a service level agreement (SLA) to an application.
 - Measured by metrics such as instructions per cycle (IPC), latency of servicing work, etc.
- Interference from unrelated workloads due to
 - Sharing of memory, caches, fabric, etc.
 - Arbitration for shared resources
- Quality of service mechanism needed - **gap**
 - Tag requests by originating workloads
 - Configure resource allocations based on tags
 - Cache capacity
 - Interconnect B/W
 - Memory B/W
 - Monitor resource usage
 - Measure and tune resource requirements
 - Enables techniques such as pseudo-locking critical code/data in caches to meet real time deadlines



Confidential Computing

- Hardware-based trusted execution environment
 - Data-in-use protection
 - Prevent unauthorized access or modification of data
- Devices in untrustworthy/unguarded locations/geographies
- Insider threats
- IP protection
- Address regulations around data processing
 - GDPR, CCPA, HIPAA, PCI, etc.
- TEE capabilities
 - Authenticated launch
 - Confidentiality, integrity, replay protection
 - Attestation
 - Sealing
 - Multi-party computation
 - Confidential computing on CPU as well as accelerators
- Virtual machines as TEE provide most flexible unit of protection
 - Cyclic executive, to RTOS, to general OS



- Standards and references
 - DMTF security protocol and data model (SPDM)
 - IETF Remote Attestation Procedures (RATS)
 - IETF Concise Reference Integrity Manifest (CoRIM, CoMID, CoSWID)
 - DMTF security protocol and data model (SPDM)
 - PCI-SIG Component Measurement and Authentication (CMA)
 - PCIe and CXL link Integrity and Data Encryption (IDE)

Other capabilities

- Cryptography
 - Encryption and hashing - AES, SHA, SM3, SM4, CLMUL
 - Entropy sources
 - Systems needing high performance hardware accelerations such as vector cryptography ISA - gap
 - Data encryption and protection with low overheads
- Low latency and Energy Efficient synchronization primitives
 - Work submission and completion signaling
 - Between applications
 - Between applications and devices/accelerators
- High resolution and synchronized Time
 - Nanosecond granularity
 - Time sensitive networking
- Telemetry - power, performance, health, etc.

Conclusions

- RISC-V provides a strong base for building secure, resilient, performant embedded systems
- Few gaps identified to complete the portfolio of foundational technologies
 - IOMMU, QoS, CFI, RAS, vector cryptography
- Calling on embedded SIG to help define the strategy to address these gaps